

Introduction to FastText

Leonardo Campillos Llanos

LIMSI - CNRS

- Developed by Facebook team
- Allows computing word vectors
- Shallow neural model relying on a hidden layer
- In the hidden layer, a sentence is represented by averaging the vector representations of each word.
- This text representation is then input to a linear classifier
→ hierarchical softmax (reduces computational complexity)

- Extension of word2vec
- Both Word2Vec architectures available:
 - Continuous bag of words (CBOW): the context is used to predict target word; does not capture the order of words
 - Skip-gram: each word is used to predict a target context

Mikolov, T., and Dean, J. (2013). Distributed representations of words and phrases and their compositionality. Advances in neural information processing systems.

- Subword information (character n-grams) can be considered:
 - Words are represented as the sum of the n-gram vectors
→ Word morphology
 - Processing of out-of-vocabulary (OOV) words
→ OOV words are represented by summing the representation of character n-grams

```
./fasttext print-vectors model.bin < OOV_words
```

- Can be used for supervised classification tasks:

```
./fasttext supervised -input train -output model
```

```
./fasttext test model.bin test.txt k
```

k: optional argument to compute precision/recall at the given value (default is 1)

```
./fasttext predict model.bin test.txt k
```

- By default, target values need to be declared with `__label__`:
This is a positive sentence `__label__1`
This is a negative sentence `__label__0`

Parameters (list not exhaustive; default value in brackets)

- cbow / skipgram
- vector dimension [100]
- context window (before and after the target word) [5]
- negative: negative sample size [5]; "negative sampling only calculates the probability with reference to a set number of other randomly chosen negative words" (Chiu et al 2016)
→ The larger it is, the slower it takes to train.
- learning rate [0.05]
- sampling threshold [0.0001]

Parameters (list not exhaustive; default value in brackets)

- minimum number of word occurrences [5]
- **word n-grams** [1]
- **minimun** [3] and **maximum** [6] **length of character-n-grams**
- **pretrained vectors** (.vec format)
- number of threads [12]
- number of buckets: n° of n-gram keys in the vocabulary hash [2 mill.]

/!\Only UTF8 encoding /!

- Python version of FastText:
<https://pypi.python.org/pypi/fasttext>

Sample use cases

- Context of modeling out-of-vocabulary terms and relate them to semantic types of in-vocabulary terms
- Hypothesis: a new term (δ) will share semantic properties of known terms (τ) occurring in similar contexts → **word-similarity task**
- Pretrained vectors on a subset (>7M tks) of the European Medicine Agency corpus

<http://opus.lingfil.uu.se/EMEA.php/>

Follow the activities prepared...

insomniantes

('insomnies', 0.8454272150993347) → disease

('insomnie3', 0.7903878688812256) → OOV (typo)

('insomniea', 0.7886084914207458) → OOV (typo)

('insomnia', 0.7879734635353088)

('insomnie', 0.7849773168563843) → disease

('somnifères', 0.7168412208557129) → OOV

('anxiétés', 0.6793801188468933) → disease

('anxiété', 0.674299418926239) → disease

('délirantes', 0.6704117059707642) → disease

('anxiété†', 0.6630647778511047) → OOV (typo)

Points to discuss?

Influence of:

- Corpus size and source
- Lemmatization and normalization (lowercase, removing hyphen and accents...)
- General / domain applications
- Tasks...



References I



Bojanowski, Piotr, Grave, Edouard, Joulin, Armand, and Mikolov, Tomas

Enriching Word Vectors with Subword Information.

arXiv preprint arXiv:1607.04606. 2016



Chiu, B., Crichton, G., Korhonen, A., and Pyysalo, S.

How to train good word embeddings for biomedical NLP.

Proc. ACL 2016, 166. 2016

<https://aclweb.org/anthology/W/W16/W16-2922.pdf>



Joulin, Armand, Grave, Edouard, Bojanowski, Piotr, and Mikolov, Tomas

Bag of tricks for efficient text classification.

arXiv preprint arXiv:1607.04606, 2016



Le, Q. V., and Mikolov, T.

Distributed Representations of Sentences and Documents

ICML. (Vol. 14, pp. 1188-1196). 2014

<http://www.jmlr.org/proceedings/papers/v32/le14.pdf>